




NAKED COROUTINES

live with Networking
Gor Nishanov • Visual C++ Team • Microsoft




DESIGN PRINCIPLES

- **Scalable** (to **billions** of concurrent coroutines)
- **Efficient** (resume and suspend operations comparable in cost to a function call overhead)
- Seamless interaction with existing facilities **with no overhead**
- **Open ended** coroutine machinery allowing library designers to develop coroutine libraries exposing various high-level semantics, such as generators, tasks, async streams and more.
- **Usable** in environments where **exceptions** are forbidden or **not available**

COROUTINES

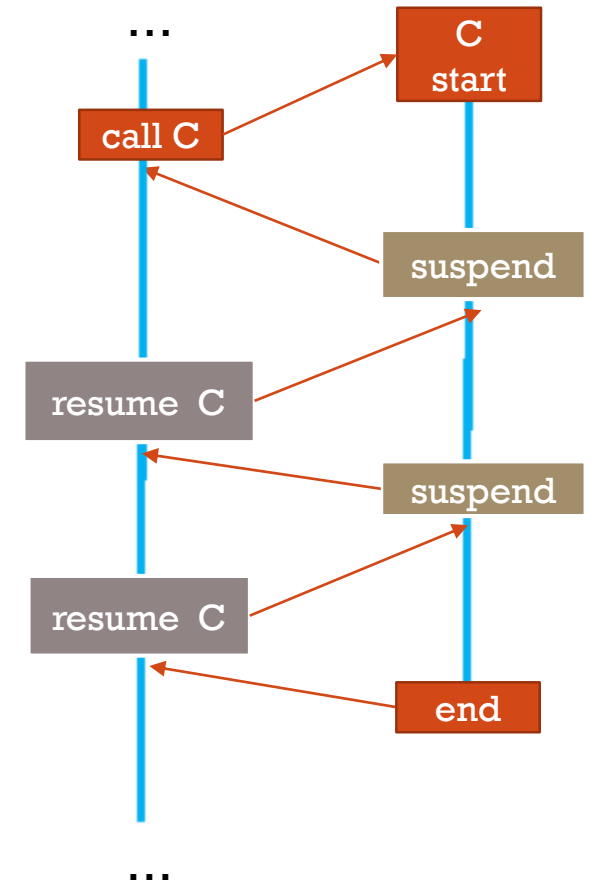
Subroutine A Subroutine B



- Introduced in 1958 by Melvin Conway
- Donald Knuth, 1968: “generalization of subroutine”

	subroutines	coroutines
call	Allocate frame, pass parameters	Allocate frame, pass parameters
return	Free frame, return result	Free frame, return eventual result
suspend	x	yes
resume	x	yes

Subroutine A Coroutine C



8.4 Function definitions

[dcl.fct.def]

8.4.4 Coroutines

[dcl.fct.def.coroutine]

Add this subclause to 8.4.

A function is a *coroutine* if it contains a *coroutine-return-statement* (6.6.3.1), an *await-expression* (5.3.8), a *yield-expression* (5.20), or a range-based `for` (6.5.4) with `co_await`.

```
generator<char> hello() {  
    for (char ch: "Hello, world\n")  
        co_yield ch;  
}  
  
int main() {  
    for (char ch : hello())  
        cout << ch;  
}
```

```
future<void> sleepy() {  
    cout << "Going to sleep...\n";  
    co_await sleep_for(1ms);  
    cout << "Woke up\n";  
    co_return 42;  
}  
  
int main() {  
    cout << sleepy.get();  
}
```

```

C++ source #1 x
63
64 template <typename T>
65 generator<T> seq() {
66     for (T i = {}; ++i)
67         co_yield i;
68 }
69
70 template <typename T>
71 generator<T> take_until(generator<T>& g, T limit) {
72     for (auto&& v: g)
73         if (v < limit) co_yield v;
74         else break;
75 }
76
77 template <typename T>
78 generator<T> multiply(generator<T>& g, T factor) {
79     for (auto&& v: g)
80         co_yield v * factor;
81 }
82
83 template <typename T>
84 generator<T> add(generator<T>& g, T adder) {
85     for (auto&& v: g)
86         co_yield v + adder;
87 }
88
89 int main() {
90     auto s = seq<int>();
91     auto t = take_until(s, 10);
92     auto m = multiply(t, 2);
93     auto a = add(m, 110);
94     return std::accumulate(a.begin(), a.end(), 0);
95 }

```

x86-64 clang 5.0.0 (Editor #1, Compiler #1) x

x86-64 clang 5.0.0 -std=c++14 -O2 -stdlib=libc++ -fcoroutines-ts

A 11010 .LX0: .text // \s+ Intel Demangle

```

1 main: # @main
2     mov     eax, 1190
3     ret

```

<https://godbolt.org/g/26viuZ>

clang version 5.0.0 (tags/RELEASE_500/final 312636)-cached

GIFTS FROM TORONTO 2017



Coroutine TS



Networking TS

OPENING THE NETWORKING TS BOX!

`io_context`

+ and more nifty things

`executors`

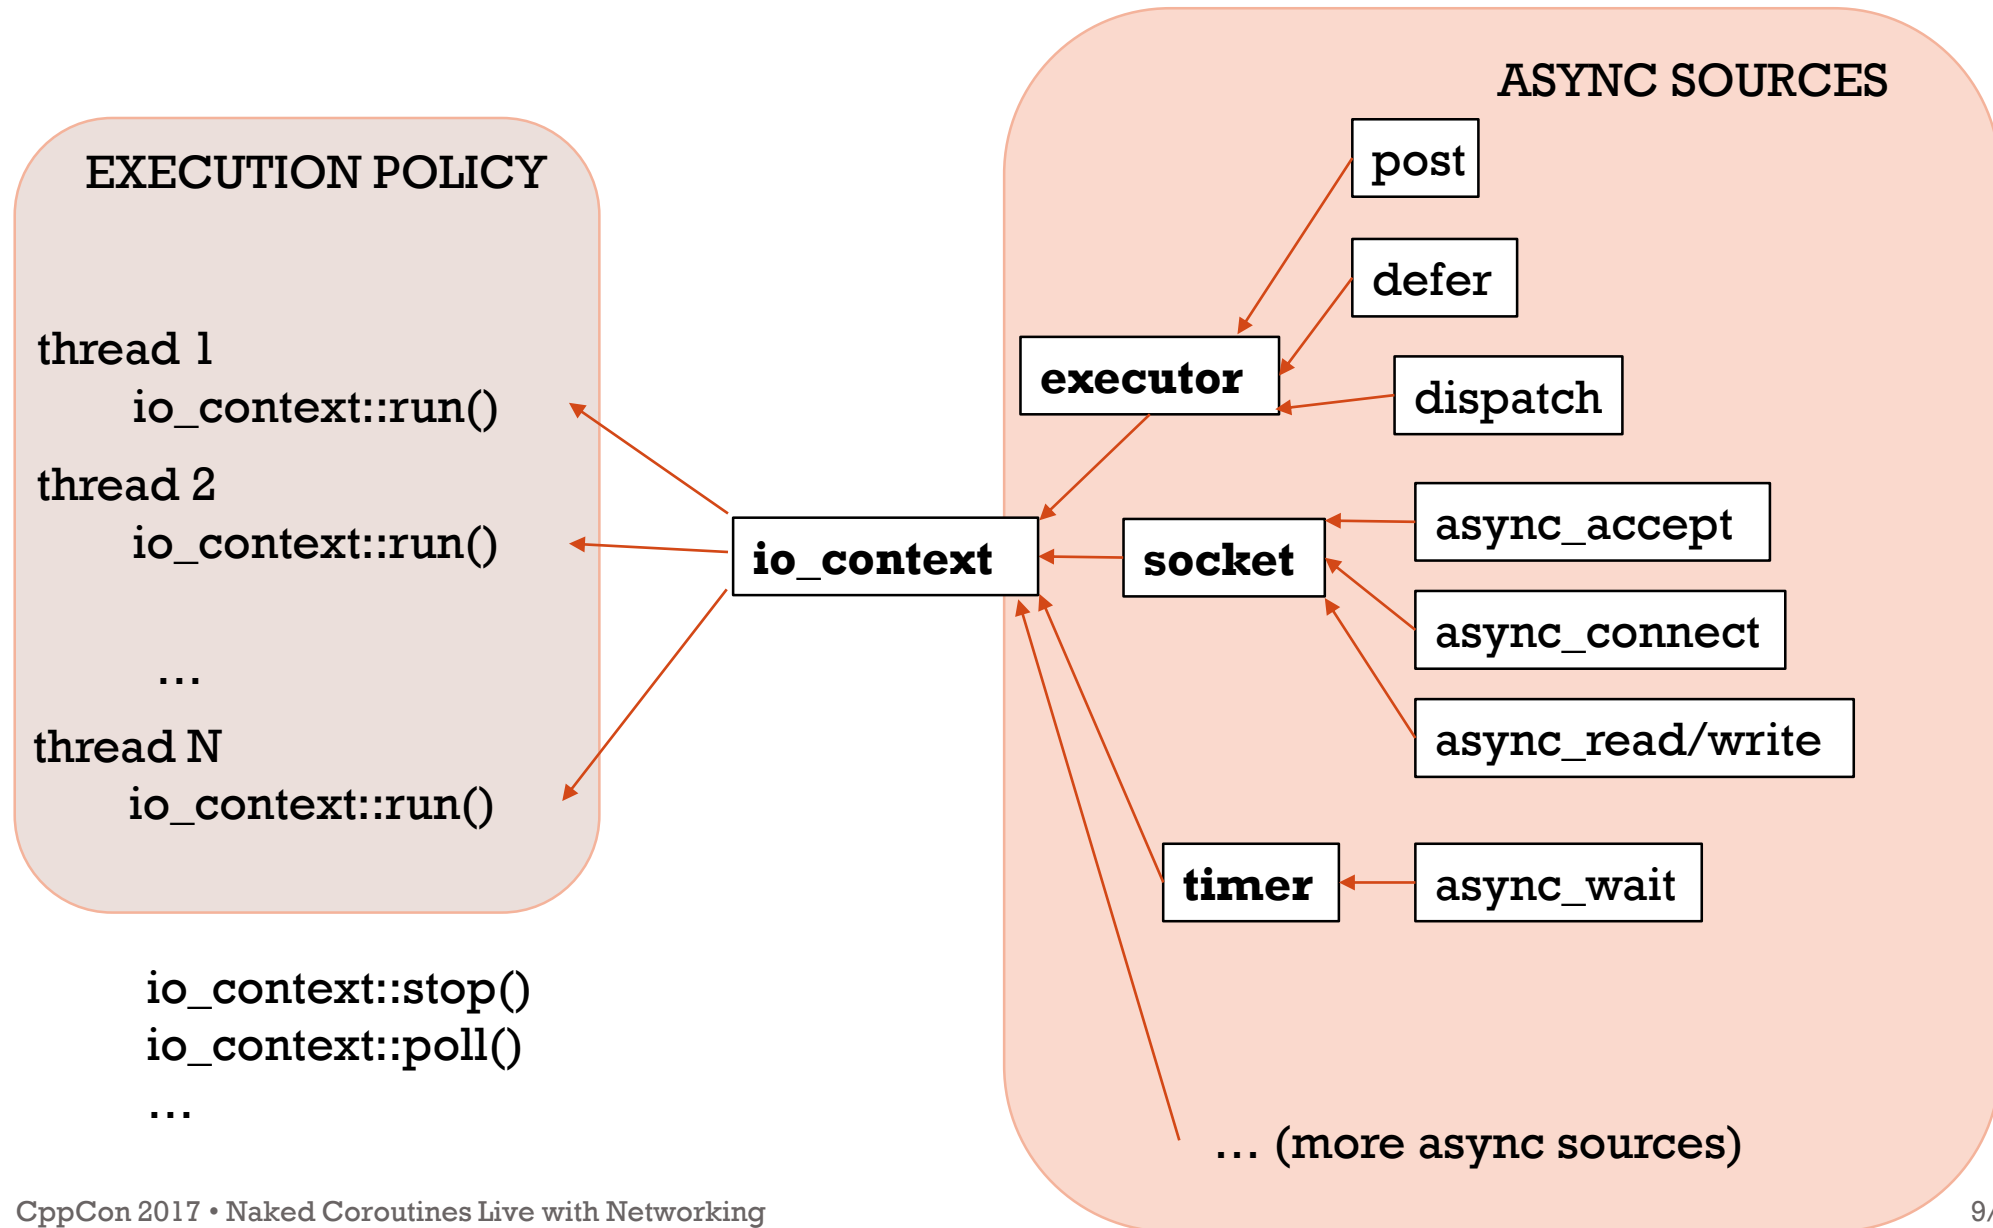
`timers`

`tcp::endpoint`
`tcp::socket`
`tcp::acceptor`
`tcp::resolver`
`tcp::iostream`



`udp::endpoint`
`udp::socket`
`udp::resolver`

NETWORKING TS – IO_CONTEXT



SIMPLE TIMER EXAMPLE

```
int main() {  
    io_context io;  
  
    system_timer slow_timer(io, hours(15));  
    slow_timer.async_wait([](auto) {  
        puts("Timer fired");  
    });  
  
    system_timer fast_timer(io, seconds(1));  
    fast_timer.async_wait([&io](auto) {  
        io.stop();  
    });  
  
    io.run();  
}
```

BEAUTIFUL TCP SERVER

```
struct session {
    session(net::io_context &ioc, net::ip::tcp::socket s, size_t block_size)
        : io_context_(ioc), socket_(std::move(s)), block_size_(block_size),
          buf_(block_size), read_data_length_(0)
    {}

    void start() {
        std::error_code set_option_err;
        net::ip::tcp::no_delay no_delay(true);
        socket_.set_option(no_delay, set_option_err);
        if (!set_option_err) {
            socket_.async_read_some( net::buffer(buf_.data(), block_size_),
                                     make_custom_alloc_handler( allocator_,
                                                                [this](auto ec, auto n) { handle_read(ec, n); }));
            return;
        }

        net::post(io_context_, [this] { destroy(this); });
    }

    void handle_read(const std::error_code &err, size_t length) {
        if (!err) {
            read_data_length_ = length;
            async_write(socket_, net::buffer(buf_.data(), read_data_length_),
                        make_custom_alloc_handler( allocator_,
                                                  [this](auto ec, auto) { handle_write(ec); }));
            return;
        }

        net::post(io_context_, [this] { destroy(this); });
    }

    void handle_write(const std::error_code &err) {
        if (!err) {
            socket_.async_read_some(net::buffer(buf_.data(), block_size_),
                                     make_custom_alloc_handler( allocator_,
                                                                [this](auto ec, auto n) { handle_read(ec, n); }));
            return;
        }

        net::post(io_context_, [this] { destroy(this); });
    }

    static void destroy(session *s) { delete s; }

private:
    net::io_context &io_context_;
    net::ip::tcp::socket socket_;
    size_t block_size_;
    std::vector<char> buf_;
    size_t read_data_length_;
    handler_allocator allocator_;
};

struct server {
    server(net::io_context &ioc, const net::ip::tcp::endpoint &endpoint,
          size_t block_size)
        : io_context_(ioc), acceptor_(ioc, endpoint), block_size_(block_size)
    {
        acceptor_.listen();
        start_accept();
    }

    void start_accept()
    {
        acceptor_.async_accept(
            [this](auto ec, auto s) { handle_accept(ec, std::move(s)); });
    }

    void handle_accept(std::error_code err, net::ip::tcp::socket s)
    {
        if (!err) {
            session *new_session = new session(io_context_, std::move(s), block_size_);
            new_session->start();
        }
        start_accept();
    }

private:
    net::io_context &io_context_;
    net::ip::tcp::acceptor acceptor_;
    size_t block_size_;
};
```

UNBOXING THE COROUTINES

and that is all you get!

`suspend_never`

`suspend_always`

`coroutine_handle`

`coroutine_traits`



`co_await`

`co_yield`

`co_return`

2)

LIVE



THE EASY WAY

ASYNCHRONOUS INITIATING FUNCTION

```
template <typename BufferSequence, typename CompletionToken>
auto async_xyz(BufferSequence const& buffers, CompletionToken handler)
{
    async_completion<CompletionToken, void(std::error_code, std::size_t)> init(handler);

    impl.real_async_xyz(buffers, init.completion_handler);
    return init.result.get();
}
```

CompletionToken →

- What to return
- What to pass as a callback to real implementation
- What executor to complete on
- What allocator to use

TRAIT SPECIALIZATION FOR USE_BOOST_FUTURE

```
template <>
struct async_result<use_boost_future_t, void(std::error_code, size_t)> {
    using return_type = boost::future<size_t>;
    struct completion_handler_type {
        boost::promise<size_t> p;
        completion_handler_type(use_boost_future_t const&) {}
        void operator() (std::error_code const& ec, size_t n) {
            if (ec) p.set_exception(std::system_error(ec));
            else p.set_value(n);
        }
    };
    explicit async_result(completion_handler_type &h) : fut(h.p.get_future()) {}
    auto get() { return std::move(fut); }
private:
    boost::future<size_t> fut;
};
```

LIVE

BEAUTIFUL TCP SERVER

```
struct session {
    session(net::io_context &ioc, net::ip::tcp::socket s, size_t block_size)
        : io_context_(ioc), socket_(std::move(s)), block_size_(block_size),
          buf_(block_size), read_data_length_(0)
    {}

    void start() {
        std::error_code set_option_err;
        net::ip::tcp::no_delay no_delay(true);
        socket_.set_option(no_delay, set_option_err);
        if (!set_option_err) {
            socket_.async_read_some( net::buffer(buf_.data(), block_size_),
                                     make_custom_alloc_handler( allocator_,
                                                                [this](auto ec, auto n) { handle_read(ec, n); }));
            return;
        }

        net::post(io_context_, [this] { destroy(this); });
    }

    void handle_read(const std::error_code &err, size_t length) {
        if (!err) {
            read_data_length_ = length;
            async_write(socket_, net::buffer(buf_.data(), read_data_length_),
                        make_custom_alloc_handler( allocator_,
                                                  [this](auto ec, auto) { handle_write(ec); }));
            return;
        }

        net::post(io_context_, [this] { destroy(this); });
    }

    void handle_write(const std::error_code &err) {
        if (!err) {
            socket_.async_read_some(net::buffer(buf_.data(), block_size_),
                                     make_custom_alloc_handler( allocator_,
                                                                [this](auto ec, auto n) { handle_read(ec, n); }));
            return;
        }

        net::post(io_context_, [this] { destroy(this); });
    }

    static void destroy(session *s) { delete s; }

private:
    net::io_context &io_context_;
    net::ip::tcp::socket socket_;
    size_t block_size_;
    std::vector<char> buf_;
    size_t read_data_length_;
    handler_allocator allocator_;
};

struct server {
    server(net::io_context &ioc, const net::ip::tcp::endpoint &endpoint,
           size_t block_size)
        : io_context_(ioc), acceptor_(ioc, endpoint), block_size_(block_size)
    {
        acceptor_.listen();
        start_accept();
    }

    void start_accept()
    {
        acceptor_.async_accept(
            [this](auto ec, auto s) { handle_accept(ec, std::move(s)); });
    }

    void handle_accept(std::error_code err, net::ip::tcp::socket s)
    {
        if (!err) {
            session *new_session = new session(io_context_, std::move(s), block_size_);
            new_session->start();
        }
        start_accept();
    }

private:
    net::io_context &io_context_;
    net::ip::tcp::acceptor acceptor_;
    size_t block_size_;
};
```


SAME BEAUTIFUL TCP SERVER BUT NOW WITH A BIGGER FONT

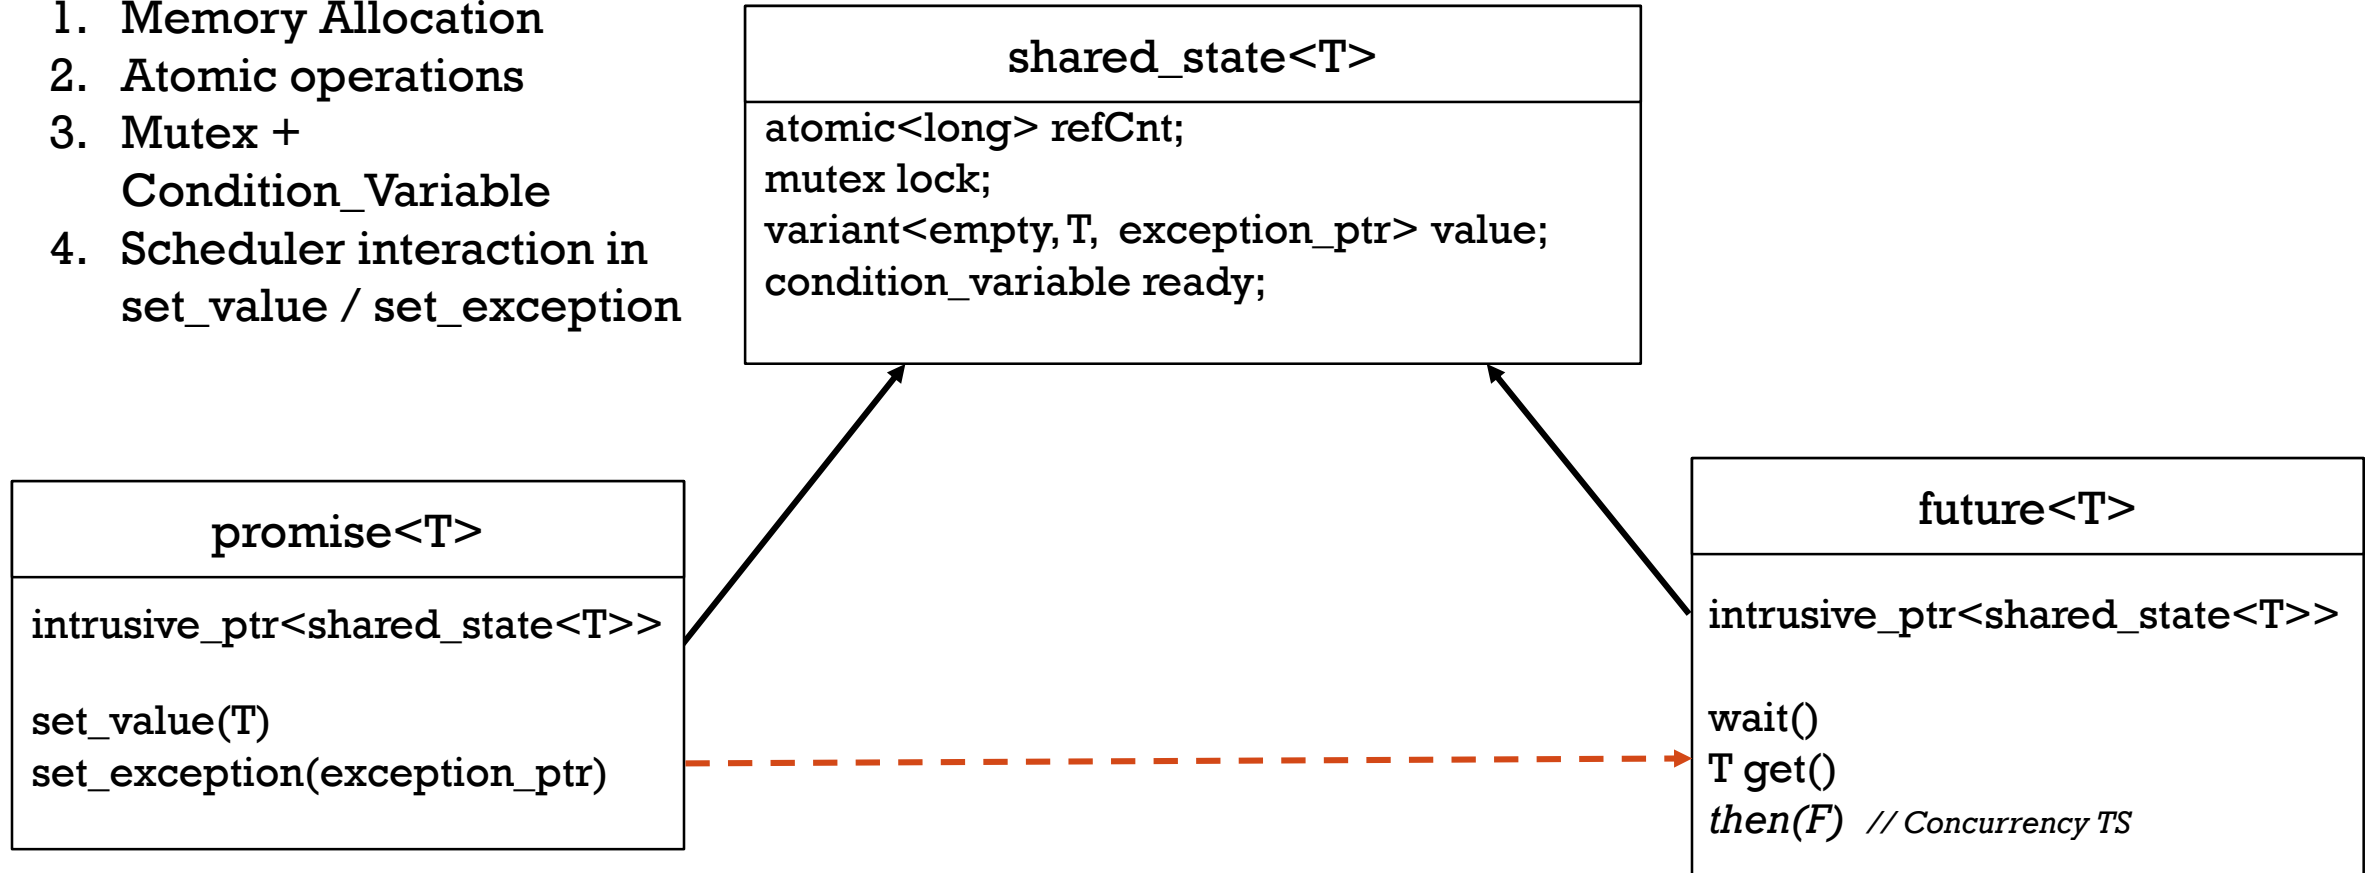
```
std::future<void> session(tcp::socket s, size_t block_size)
{
    s.set_option(tcp::no_delay(true));
    std::vector<char> buf(block_size);

    for(;;) {
        size_t n = co_await async_read_some(s, buffer(buf.data(), block_size));
        n = co_await async_write(s, buffer(buf.data(), n));
    }
}
```

```
std::future<void> server(io_context &io, tcp::endpoint const &endpoint,
                        size_t block_size)
{
    tcp::acceptor acceptor(io, endpoint);
    acceptor.listen();
    for (;;)
        session(co_await async_accept(acceptor), block_size);
}
```

STD::FUTURE<T> AND STD::PROMISE<T>

1. Memory Allocation
2. Atomic operations
3. Mutex +
Condition_Variable
4. Scheduler interaction in
set_value / set_exception





COMPLICATIONS

Cancellation and allocation

LIVE



BEYOND THE TS

Two possible additions to C++ Coroutines

SYMMETRIC CONTROL TRANSFER

TWEAK FINAL SUSPEND

```
template <typename T> struct task {
    struct promise_type {
        variant<monostate, T> result;
        coroutine_handle<> waiter;

    _
    auto final_suspend() {
        struct Awaiter {
            promise_type* me;
            bool await_ready() { return false; }
            void await_suspend(coroutine_handle<>) {
                me->waiter.resume();
            }
            void await_resume() {}
        };
        return Awaiter{this};
    };
    template <typename U> void return_value(U &&value) {
        result.emplace<1>(std::forward<U>(value));
    }
};
```

1. Memory Allocation
2. Atomic operations
3. Mutex + Conditional Variable
4. Scheduler interaction in set_value / set_exception

Tail Call

GOR NISHANOV

C++ Coroutines Under The Cover

- Available only in clang trunk
- Not in MSVC or clang 5
- Not part of the TS (yet)

```
coroutine_handle<> await_suspend(coroutine_handle<>) {
    return me->waiter;
}
```

PEEKING AT COROUTINE ARGUMENTS FROM PROMISE

```
// Coroutine object returned in an usual place
HRESULT f(X x, Y y, Z z, SomeSmartPtr<MyCoro>* p);

// Would like have access to executor in initial_suspend
void g(executor& e);

// Would like to check whether we are cancelled at every
// suspend point
void h(cancellation_token& c);
```

NOT VERY GOOD WORKAROUND

```
struct promise_type {  
  
    template <typename... Args>  
    void* operator new(size_t size, Args const&... args) {  
        // stash what you need into a thread_local  
    }  
  
    promise_type() {  
        // get what you wanted out of a thread_local  
    }  
    ...  
};
```

LET PROMISE CONSTRUCTOR PEEK AT ARGS!

```
struct promise_type {  
  
    template <typename... Args>  
    promise_type(Args const&... args) {  
        // get what you want  
    }  
    ...  
};
```

- Opt-in feature. Empty construct will work fine
- Will observe stable parameters (parameter copies)
- Implicit object parameter passed as a first argument
- Not part of the TS
- Not available in any compiler

CONCLUSION

- Networking and Coroutine TS are great together
- At the moment, for the best performance use “the hard way”
- Hopefully can be addressed before C++20 ships
- Coroutines are available in
 - MSVC 2017 (/await)
 - clang 5.0 (-fcoroutines-ts -stdlib=libc++)
- Networking TS implementation:
 - <https://github.com/chriskohlhoff/networking-ts-impl>
- Look at good open source coroutine libraries:
Example: <https://github.com/lewissbaker/cppcoro>
- Snippets we used during the live part will be available at:
https://github.com/GorNishanov/await/tree/master/2017_CppCon



QUESTIONS?



BACKUP

C++ COROUTINES

COMMON PATTERN FOR ASYNC AND SYNC I/O


async

```
future<int> tcp_reader(int total)
{
    char buf[64 * 1024];
    auto conn = co_await Tcp::Connect("127.0.0.1", 1337);
    for (;;)
    {
        auto bytesRead = co_await conn.read(buf, sizeof(buf));
        total -= bytesRead;
        if (total <= 0 || bytesRead == 0) co_return total;
    }
}
```

sync

```
expected<int> tcp_reader(int total)
{
    char buf[64 * 1024];
    auto conn = co_await Tcp::Connect("127.0.0.1", 1337, block);
    for (;;)
    {
        auto bytesRead = co_await conn.read(buf, sizeof(buf), block);
        total -= bytesRead;
        if (total <= 0 || bytesRead == 0) co_return total;
    }
}
```

SNEAK PEEK
at the bright
future



C++ COROUTINES

REACTIVE STREAMS MEET COROUTINES

Source

Produces 0.1.2.3...
each 1ms

```
async_stream<int> Ticks()  
{  
    for (int tick = 0;; ++tick)  
    {  
        co_yield tick;  
        co_await sleep_for(1ms);  
    }  
}
```

Transformer

Transforms stream of v_1, v_2, v_3, \dots
into a stream of
 $(v_1, t_1), (v_2, t_2), (v_3, t_3), \dots$
where t_i is a timestamp of when
 v_i was received

```
template<class T>  
auto AddTimestamp(AsyncStream<T>& S)  
{  
    for await(auto&& v: S)  
        co_yield make_pair(v, system_clock::now());  
}
```

Sink

Reduces an asynchronous
stream to a sum of its values

```
future<int> Sum(AsyncStream<int>& input)  
{  
    int sum = 0;  
    for co_await(v: input)  
        sum += v;  
    co_return sum;  
}
```

SNEAK PEEK
at the bright
future

N4134: RESUMABLE FUNCTIONS V2

GENERATORS AND ITERABLES AND AGGREGATE INITIALIZATION

generators

```
generator<char> hello() {  
    for (ch: "Hello, world\n") yield ch;  
}  
int main() {  
    for (ch : hello()) cout << ch;  
}
```

SNEAK PEEK
(more later)

Not in
N4134

constexpr
generators

```
int a[] = { []{ for(int x = 0; x < 10; ++x) yield x*x; } };
```


Equivalent to

```
int a[] = { 0,1,4,9,16,25,36,49,64,81 };
```

Recursive
Generators

Checks if in-order
depth first
traversal of two
trees yields the
same sequence of
values

```
auto flatten(node* n) -> generator<decltype(n->value)> {  
    if (n != nullptr) {  
        yield flatten(n->left);  
        yield n->value;  
        yield flatten(n->right);  
    }  
}  
bool same_fringe(node* tree1, node* tree2)  
{  
    auto seq1 = flatten(tree1);  
    auto seq2 = flatten(tree2);  
    return equal(begin(seq1), end(seq1),  
                 begin(seq2), end(seq2));  
}
```



WARNING: DO NOT DO THIS!

```
template <typename R>
auto operator co_await(std::future<R>&& f) {
    struct Awaiter {
        std::future<R>&& f;
        bool await_ready() {
            return f.wait_for(0s) != future_status::ready();
        }
        void await_suspend(std::experimental::coroutine_handle<> h) {
            std::thread([&]() mutable {f.wait(); h.resume();}).detach();
        }
        auto await_resume() { return f.get(); }
    };
    return Awaiter{ std::forward<std::future<R>>(f) };
}
```


WITH USE_FUTURE, BEHAVES AS IF IT WERE

```
template <typename BufferSequence>
auto async_xyz(BufferSequence const& buffers, use_future_t)
{
    std::promise<size_t> p;
    auto result = p.get_future();

    impl.real_async_xyz(buffers, [p = std::move(p)](auto ec, auto n) {
        if (ec) p.set_exception(std::system_error(ec));
        else (p.
    ));
    return result;
}
```

Implementation:

https://github.com/chriskohlhoff/networking-ts-impl/blob/master/include/experimental/_net_ts/impl/use_future.hpp